

The GPU Showdown — Instructor Key

1. $15.35/5.79 = 2.65\times$ approximately. It is an observed whole-system ratio; CPU, OS, backend, driver, load, and residency differ.
2. April, Maise, Faith, Brandy, STF125, STF259, Grace, Morgan. The ranking does not establish architectural causation; STF125 has six explicit timeouts and STF259's 15.92-second all-six-model value is not the ranking metric.
3. A 3 GB limit may force offload or restrict model choice. The source supports a residency caveat, not a measured causal multiplier.
4. Examples: CPU/host and backend; control by matching host/software and recording residency.
5. Defensible answer: "In these recorded runs, Maise's observed latency was lower than both STF125's 5060 Ti and STF259's 3060 Ti; uncontrolled host/software/residency differences prevent attributing that difference to GPU family."
6. It prevents a false benchmark row. A gate outcome has no valid latency value.
7. Same model, prompts, repetitions, software, host conditions, residency verification, and explicit timeout/error accounting.